

MVP End-to-End Master Command

Thirteen front-end commands to make the version we already have work, end to end, for a real customer — plus the backend blockers that no amount of front-end work can move.

Field	Value
Prepared	19 August 2026 . Every figure marked <i>measured</i> was taken on this date against <code>api.servana.com.ph</code> or the repository itself.
Repository	<code>C:\Users\paulg\OneDrive\Desktop\servana_client-unified</code> — the unified working copy. Supersedes <code>servana_client-mobile</code> .
Remote	<code>github.com/Upupapp/ServanaClientAPP</code> . Nothing in this document has been pushed.
Gates today	All green — <code>dart format</code> exit 0 (732 files, 0 changed); <code>flutter analyze</code> No issues found ; <code>flutter test</code> 2590 passed , 6 skipped. Backend: <code>npm run verify</code> exit 0, 6509 tests .
Revision	Rev 4 , 20 Aug. TABs 01, 02, 04, 05 closed . Backend B1, B2, B7, B9 closed ; production healthy for 11 hours. A full front-end sweep has since fixed 34 layout defects and taken rendered screen coverage from 11 to 56 of 62 .
Scope	Front-end only . Backend items appear as blockers and evidence, never as tasks in a TAB.
Definition of MVP	One customer can install the app, create an account, sign in, find a service, book it, pay for it, watch it happen, talk to the provider, and review it — without hitting a screen that lies, dead-ends, or cannot be completed.

Resolved since Rev 1. Production was answering `500` on every database-backed route when this document was written. It is now healthy — `/readyz` reports `ready:true`, all five dependencies up, and the catalog serves **95 services**. The cause was **not** what it first appeared: see **Section B**, which is now a post-mortem rather than a blocker list. **Every front-end TAB below is now verifiable.**

Front-end commands are ordered so that each one is verifiable the moment the backend is healthy. TAB 01 exists to make “works end to end” a measurement rather than an opinion.

Section A — What is already true

CLOSED TODAY, IN THIS REPOSITORY

Commit	What
e29fcd9	Unified two divergent <code>main</code> branches — 31 production-launch commits on GitHub and 43 front-end SWEEP commits locally that had never seen each other.
00320db	Two login-failure assertions that could not fail (the analyzer rejected them as tautologies) now classify through the base type, as production does.
1bd9e39	Analyzer taken back to zero issues across the merged tree.
7c4a178	The app no longer tells customers their password is wrong when the server is down. See TAB 02.
cd4a57b	A stale notifications list said nothing, and the icon invented a cause. TAB 02 closed.
dd6d988	Withdrew the affordances a customer cannot finish. TAB 04 closed.
ff63a30	A 500 drew the wifi-off icon in Search. TAB 05 closed.

One theme runs through all four. Every defect was the UI asserting a cause it did not know — a wrong password, a broken network, a fresh list. None was a crash, and none would have been caught by a test that did not ask what the customer is being told.

FRONT-END AREAS MEASURED HEALTHY — DO NOT SPEND MVP TIME HERE

- **Token refresh.** Single-flight (`_refreshInFlight`), one owner for sign-out, 401 handled in one place. No logout loop of the kind the provider portal has.
- **Catalog to booking handoff.** `canonical_booking_handoff.dart` routes a canonical Service into a real checkout, and an **unknown category falls to the generic flow rather than failing** — a new category is bookable the day an admin creates it, with no app release.
- **Booking submission error copy.** Substring-based, but its default is `serverFailure` and reads “Something went wrong. Please try again.” It fails safe; it never blames the customer.
- **Route reachability.** 67 `GoRoute` s, derived from screen constants rather than string literals.
- **Capability gating.** Deny-by-default, build-time only, not remote-configurable. All 15 capabilities are OFF, so the shipped app is 100% legacy transport. **For MVP this is correct — leave it alone** (TAB 13).

Section B — Backend blockers (not front-end work)

These are recorded because they determine whether the front end can be verified at all. **No TAB in Section C asks anyone to fix these.**

B1 — CLOSED: THE OUTAGE, AND WHAT ACTUALLY CAUSED IT

Production is healthy. `/readyz` is `ready:true` with all five dependencies up; `/api/catalog` returns **200** with 3 categories, 12 subcategories and **95 services**.

The first diagnosis was wrong in an instructive way. Every database-backed route was answering 500 with `SASL: SCRAM-SERVER-FIRST-MESSAGE: client password must be a string` — the `pg` driver being handed `undefined`. That looked like one broken credential. It was a symptom of three separate faults stacked on each other.

THE REAL CHAIN, MEASURED

#	Fault	Evidence
1	Production ran out of the CI runner's scratch directory. PM2's script path was <code>/home/github-runner/actions-runner/_work/servana_api/servana_api/dist/app.js</code> .	A self-hosted runner service was active on the production box .
2	A deploy workflow wiped the running build. Step 3 Checkout code succeeds — and deletes <code>dist/</code> in the directory production executes from — then step 8 fails, so steps 12 Build and 15 Restart PM2 never run.	Runs at 10:56, 11:14, 13:58, 14:05, 14:08 all failed . Three during the session.
3	The failing step was an OOM, not a type error. <code>typecheck</code> passes; <code>typecheck:tests</code> aborts with <code>JavaScript heap out of memory, exit 134</code> .	The box has 961 MB RAM and 1 CPU .

A live Node process keeps its code in memory after the files are deleted, so the app kept serving for hours on a `dist/` that no longer existed. The restart is what exposed it — as a hard **502 on everything, including** `/healthz`. And the DB password was missing because step 15 never ran, so the process had never been restarted with it.

The lesson worth keeping: a memory-heavy typecheck, running on the production host, in the directory production executes from, could take the whole platform down by running out of RAM. Each layer was individually defensible. Stacked, they made a failing test a customer-facing outage.

WHAT WAS DONE

- Production moved to** `/var/www/servana_api` — outside any directory CI owns. A checkout can no longer reach it. PM2 restarted from there, `pm2 save` written.
- Actions disabled on** `Upupapp/servana_api` (`enabled:false`, verified). Note these were *self-hosted* runs, which consume no GitHub Actions minutes — they were stopped because they were destroying production, which is the better reason.
- Env file repaired.** Two lines were malformed for shell sourcing: `PAYMONGO_LINK = ...` (spaces around `=` runs the name as a command) and `PAYMONGO_WEBHOOK_SECRET= ...` (a per-command assignment that does not persist). Both variables were therefore **absent** from the process — including the secret that verifies payment callbacks. Normalised, backed up; the process now holds **7 of 7** PayMongo variables.
- `/etc/servana_api.env` installed (`root:root 600`) so nothing depends on the runner's home directory.

B2 — CLOSED: THE CATALOG ROUTE SHADOWING IS FIXED AND DEPLOYED

`GET /api/catalog` answered **401** on 18 August because a booking wildcard was mounted ahead of it. It now returns **200** with a 49 KB payload, so the mount-order fix **is** deployed. TAB 05 is unblocked and closed.

Confirmed while checking: the live payload carries `"slug":"home-maintenance"` — hyphens. Home's curated keys use underscores (`beauty_wellness`). The trap TAB 05 warns about is real; route by `catalog_categories.id`, never the slug.

B3 — COVERAGE AND BRANCH SELECTION ARE KEYED TO LEGACY FAMILY IDS

GET /api/services/:familyId/branches and /coverage-geo take a legacy family id. A canonical Service created through the Admin API has no legacy option, so neither is reachable for it. The front end can only degrade honestly here (TAB 07); making them work is backend scope.

B4 — THERE IS NO CANONICAL BOOKING CREATE

POST /api/bookings is classified KEEP with no /api/v1 successor and none planned. This is why the capability enum has bookingReads and bookingLifecycle but deliberately no booking. It does not block MVP — the legacy write is what ships — but it does block any future claim that the booking domain has migrated.

B5 — VERIFYSMOBILE NEEDS A FIREBASE PHONE CREDENTIAL

The canonical endpoint takes {idToken}, not {mobileNumber, otp}. That is a redesign, not a rename. It is why mobile-number login is still “coming soon” in the UI (TAB 04).

B7 — CLOSED: DOTENV RAN AFTER THE DB POOL WAS BUILT

The cause was worse than “load order”. config.ts — imported by 53 modules, and the module that captures the database credentials into a module-level const — guarded its own load with if (!process.env.NODE_ENV). NODE_ENV is always set in production, so it never ran; and its path “./.env” resolved against the process CWD, so it pointed at the wrong file even when it could have fired. Combined with app.ts calling dotenv.config() after its imports, .env had never supplied the database credentials in production at all. The API only ever worked because PM2 happened to carry them in its environment — which is exactly why a restart that lost them took the platform down.

Fixed and proven on production via src/env/loadEnv.ts, imported first in config.ts and app.ts. With the variables explicitly unset in the shell:

```
env -u DB_PASSWORD -u DB_USER node -e "require('./dist/config')"
-> injected env (28) from .env
-> password resolved: true    password is a string: string
```

That last line is the exact inverse of the outage error. Production then survived a plain pm2 restart with no set -a prefix. The prefix is now a belt, not the braces.

B8 — SERVANA.LOCATIONS IS READ BUT DOES NOT EXIST

Surfaced by the project's own npm run db:diagnose. adminCreateBookingService.ts:430 500s whenever a serviceLocationId is supplied. Independent of the outage.

B9 — CLOSED: TYPECHECK:TESTS RUNS ON A DEV MACHINE

It exhausts the heap on a 961 MB server (exit 134) and passes exit 0 on a developer machine. Decision recorded in the backend's docs/DEPLOY_AND_GATE_POLICY.md.

Three gates could not pass on the machine that is now the only gate. With CI off, scripts/hooks/pre-push runs npm run verify, so a Linux-only assumption blocks every push — and the obvious workaround is --no-verify, which removes the gate rather than repairing it.

deploy-gating asserted the hook's exec bit via statSync().mode & 0o111; NTFS carries no POSIX exec bit, so Node reports 666. It now asserts git ls-files -s = 100755, the mode that actually travels. **vacuous-ratchet** built keys with test.path.replace(process.cwd() + '/', '') — backslashes against a forward slash, so the replace matched nothing and it reported “24 no longer on the list” when nothing had regressed; after normalising, 3, the true number. Both fixed and pushed.

B6 — DISPUTES ARE ADMIN-ONLY

The only legacy dispute route is admin-scoped, so `canOpenDispute` is false for every customer. The controller already exposes the capability; the screen appears the day a transport can serve one. Out of MVP scope.

Section C — The front-end commands

One TAB per command. Each is independently verifiable. **Every TAB ends with evidence, not an assertion.**

TAB	Command	Blocked by	MVP
01	Make “works end to end” a measurement	—	DONE
02	Finish the sign-in and registration honesty pass	—	DONE
03	Prove the account-creation loop closes	OTP in a human inbox	In progress
04	Remove every affordance that cannot work	—	DONE
05	Home and Search must say what is actually wrong	—	DONE
06	Booking creation across every production category	— <i>unblocked</i>	Required — next
07	Address, schedule and coverage degrade honestly	B3	Required
08	Payment: cash and PayMongo, including the unhappy paths	— <i>unblocked</i>	Required
09	My Bookings, detail, tracking and OTP completion	— <i>unblocked</i>	Required
10	Messaging the provider	— <i>unblocked</i>	Required
11	Review after completion	— <i>unblocked</i>	Required
12	Session resilience and the offline banner truthfulness	—	Required
13	Cut the MVP build with capabilities OFF	01–12	Required

TAB 01 — Make “works end to end” a measurement

WHY

“The app works” is currently an opinion held by whoever last tapped through it. The suite has 2161 tests and none of them walks a customer from install to review, because every one of them stops at a seam. That is the right design for unit tests and the wrong basis for an MVP go/no-go.

DO

- Write the ten-step customer journey down as an explicit, ordered checklist — install, create account, verify, sign in, browse, open a service, book, pay, track, message, review.
- Build an `integration_test/` walkthrough that drives as much of it as a device can without a live provider on the other end. The existing `test/support/screen_test_container.dart` already registers the real DI graph over a `MockClient`; extend that pattern rather than inventing a second one.
- For the steps a test cannot reach (a provider accepting, an OTP handed over in person), write the manual script with the exact account, exact service and exact expected screen at each step.

DONE WHEN

A single command produces a pass or fail for the automatable portion, and a one-page manual script exists for the rest, with a named test account. **Record which steps are automated and which are manual** — a walkthrough that silently skips payment reads as a green light for payment.

TAB 02 — Finish the sign-in and registration honesty pass

Largely closed today in 7c4a178 . This TAB is the remainder.

WHAT WAS WRONG, AND WHY IT MATTERED

`HttpBackend` turns an unparseable response body into `Login failed (502)`. and a JSON body with no message into `Login failed.` — and `ErrorMessageMapper` classified login failures by searching those strings for keywords. Neither `502` nor `504` appears in any list; only `500` and `503` do. So both strings fell through every branch onto the default:

```
"The email or password is incorrect."
```

An nginx gateway error in front of a restarting API — **exactly today's production condition** — told the customer their password was wrong. A customer who reads that changes a password that was never wrong, and support cannot reproduce it.

The status code is now threaded through `Backend.authenticate` and `registerCustomer` and consulted **first**. `401` and `403` still defer to the body, because the difference between an unverified account and a disabled one lives there and both are `401`.

Registration had the mirror fault: `registration_bloc` put the raw backend string on screen and showed `e.toString()` for an exception, while `ErrorMessageMapper.forRegistration` — written for exactly that, and carrying the rule *no raw backend strings or stack traces should ever reach the UI* — had **no callers at all**.

22 tests were added and **mutation-checked**: with the status branch disabled, 15 of them fail.

REMAINING WORK

- Apply the same status-first rule to `resendVerificationEmail`, which still returns a bare record with no status and sits on the verification loop TAB 03 depends on.
- Add a **widget-level** test that the sign-in screen renders the mapped copy. The current tests pin the mapper and the bloc; nothing asserts what the customer actually sees.
- Audit the four other sites that render connectivity copy without consulting `isTransport`: `address_repository.dart:127`, `booking_submission_result.dart:124`, and `notifications_screen.dart:154,196` (a wifi-off icon). A `500` is not a wifi problem.

TAB 03 — Prove the account-creation loop closes

WHY

Sign-up is the only door into the app, and it is a loop, not a step: register, mail lands, OTP or link, verified, sign in. Any break makes the app unusable for a brand-new customer, which is the only kind of customer an MVP has.

DO

- Walk the whole loop with a fresh address. Confirm the verification mail actually arrives, and confirm resend works and is rate-limited sanely.
- Confirm the app bodies match what the deployed handler accepts. The v1 request schemas are **not enforced at runtime** — `src/api/v1/register.ts` runs no validator, so the handler is the authority and it accepts legacy aliases. Do not read the schema and assume.
- Confirm a customer who abandons mid-verification can get back in rather than being stranded between states.

DONE WHEN

A brand-new address reaches a signed-in Home screen, and the exact bodies sent at each step are recorded in the TAB.

TAB 04 — Remove every affordance that cannot work

WHY

An MVP is judged by what a customer can finish. A control that is visible and cannot work costs more trust than a feature that is absent — the customer does not know it is unbuilt, so they read it as broken.

KNOWN, MEASURED

Affordance	State	MVP call
Mobile-number login	Says “coming soon”. Cannot work: the canonical endpoint needs a Firebase phone credential (B5).	Remove for MVP , ship email-only
Rewards screen	Placeholder — “Rewards coming soon”.	Remove from the drawer
Favourites	Placeholder, same pattern.	Remove from the drawer
Disputes	No screen. <code>canOpenDispute</code> false everywhere (B6).	Correctly absent — leave

DO

Sweep every screen for a control that cannot complete, and either remove it or make the disabled state say *why*. Then add a test that asserts the property rather than the instances — a test written for one orphan tends to find the next one immediately.

TAB 05 — Home and Search must say what is actually wrong

WHY

The catalog is 500ing (B1 and B2), so this is not hypothetical: it is what every customer sees today. The Home “More services” row draws nothing and Search returns nothing.

The fail-safe is working as designed — the curated four cards still render and nothing crashes. What must be true is that the customer is **told the truth**: the service listing is unavailable, not “check your connection”.

DO

- Confirm the honesty rule from TAB 02 reaches Home and Search, not just sign-in. `home_composition.dart:128` already consults `isTransport`; verify every other failure surface does.
- Confirm an absent section and a failed section stay visibly different. There is nothing to retry on an absent one, so a retry button there is a lie.
- **Do not** feed catalog categories into the Home curated grid. Home keys use underscores (`beauty_wellness`) and drive both navigation and the campaign registry; the backend slugify produces hyphens. A campaign registered against a key the catalog does not use returns null: no popup, no exception, tests pass. The catalog must only *add* what the curated four do not cover, routed by `catalog_categories.id`.

TAB 06 — Booking creation across every production category

WHY

This is the transaction. Everything before it is browsing.

The handoff resolves a canonical Service into one of two checkouts by category id (2 goes to aircon, everything else to the generic Beauty and Wellness flow). That default is deliberate and good. It has not been proven for all three production categories.

DO

- Book one service in each production category — 1 Home Maintenance, 2 Home Services, 3 Personal Care.
- Book a Service **created through the Admin API**, which has no legacy option row. This is the case where `services.id == service_options.id` stops holding, and it is the one that will break first in production.
- Confirm the payload `serviceOptionId` is the canonical `services.id` the customer actually selected, and that add-on selection does not follow the customer to the next service.

TAB 07 — Address, schedule and coverage degrade honestly

WHY

Branch and slot selection and coverage-geo are keyed to legacy family ids and are **unreachable for a canonical Service** (B3). The front end cannot fix that. What it must not do is present a coverage answer it did not get.

DO

- Decide and implement the MVP behaviour when coverage cannot be determined. An undeterminable coverage answer should **deny or ask**, never silently allow — a booking accepted outside the service area costs a real dispatch.
- Confirm saved addresses work: add, select, make primary.
- Confirm the Manila coordinate fallback is not silently applied to a booking whose real coordinates are simply absent from the response.

TAB 08 — Payment: cash and PayMongo, including the unhappy paths

WHY

Payment is where an MVP either earns money or loses trust. The client sends `CASH` or `PAYMONGO` and the backend types the column `CASH|GCASH|PAYMONGO`.

DO

- Complete a cash booking end to end.
- Complete a PayMongo booking through `PaymentWebViewScreen`, then deliberately **cancel** one and **abandon** one (close the sheet mid-flow).
- Confirm each of those three outcomes leaves a booking in a state the customer can act on. An abandoned payment that leaves an invisible half-booking is worse than a failed one.
- Confirm the amount shown at checkout is the amount charged, and that a price arriving as a JSON string rather than a number still renders. Postgres numeric arrives as int, double or string depending on value and driver.

TAB 09 — My Bookings, detail, tracking and OTP completion

DO

- Confirm a booking created in TAB 06 appears in My Bookings immediately, with the right total — not ₱0.00.
- Confirm booking detail reads through `BookingRepository` and shows status, schedule, address, provider and total.
- Confirm tracking renders when the provider is en route, and degrades honestly when there is no location rather than dropping a pin on a default coordinate.
- Confirm the completion OTP path works from the customer side.

Watch for: the domain model was once lossier than the screen own parsing — a blind swap onto the repository regressed five behaviours including the total. Pin the model against the screen before trusting it.

TAB 10 — Messaging the provider

DO

- Open the conversation from a booking, send and receive, confirm unread counts and read receipts settle.
- Confirm the conversation is reachable from both the booking detail and the inbox, and that both resolve to the same thread.
- Confirm a message that fails to send is visibly unsent and retryable, rather than silently dropped.

TAB 11 — Review after completion

DO

- Confirm review eligibility appears only after completion, and that submitting once closes the affordance.
- Confirm the rating reaches the provider aggregate the app reads back.

Note: only 4 of 9 review calls have a canonical successor, which is why `reviews` is absent from the capability enum. Irrelevant for MVP — the legacy path is what ships.

TAB 12 — Session resilience and the offline banner truthfulness

DO

- Confirm a token expiring mid-session refreshes once, not once per in-flight request, and that a dead refresh token signs the customer out exactly once.
- Confirm the offline banner appears only when the device is genuinely offline. During B1 the device is online and the server is broken — the banner must not claim otherwise.
- Confirm backgrounding the app for an hour and returning does not strand the customer on a blank screen.

TAB 13 — Cut the MVP build with capabilities OFF

WHY OFF IS THE RIGHT CALL

`/api/v1` is deployed — measured today, `GET /api/v1/catalog` returns 500 rather than 404, which means the route exists and reached the database. So the gate original premise (“v1 is unpushed”) is now out of date.

That is **not** a reason to enable it for MVP. Both transports fail identically today, so switching buys nothing, and every capability flipped on is a surface that has never carried real customer traffic. **Ship the transport that ships.** Migrating is a post-MVP command with its own evidence.

DO

- Cut a release build with all 15 capabilities off, and assert that in a test rather than trusting the default.
- Run the built artefact, not the debug one. R8, signing and the version floors are already done; confirm the shrunk build still runs the journey from TAB 01.
- Confirm the version gate can retire this build later, and that the deep links resolve.

Section F — The layout sweep (20 August)

34 defects found and fixed, every one of them **silent in a release build** — the content is clipped, no error is reported, and the customer simply sees less than was written. Rendered screen coverage went from **11 to 56 of 62**; the suite from 1,901 to **2,590** tests.

THREE ROOT CAUSES ACCOUNT FOR NEARLY ALL OF THEM

Pattern	Count	Worst instance
A bare <code>Text</code> inside a <code>Row</code>	~20	"Continue with Facebook" on the sign-in screen — 122px over at text scale 1.0 , the default
A fixed-size box or <code>Column</code> around scaling content	~8	The booking calendar: every one of its 27+ day cells overflowed at 1.3
Only ONE side of a row bounded	3	The shared <code>settings_tile</code> — already had <code>Expanded</code> on the title, so it read as handled

The third is the one that hides: somebody has already written `Expanded`, so the row looks fixed, but the value beside it grows too. **If both sides of a row can grow, both have to be able to give.**

SIX FIXES WERE IN SHARED WIDGETS

Two separate primary-button implementations, the QR placeholder on both confirmation screens, the analytics consent sheet, the booking dialog sheet (which alone carried **nine** overflows), and the settings tile. Those reach screens the matrix has not even rendered yet.

ONE DEFECT WAS NOT A LAYOUT PROBLEM AT ALL

`CreateAccountScreen` started an `overlay_tooltip` tour in `initState`, registered no tooltip item, and the app has no `OverlayTooltipScaffold` anywhere — so it **threw 500ms after every new customer opened the sign-up screen**, and the tour never appeared. The widget written to render it has no callers.

It surfaced because the screen could not be added to the matrix. I had recorded that as "a harness gap, not a layout defect". It was both: the harness could not build the screen *because the screen throws in production too*. Taking the first explanation at face value would have left a live defect hidden behind a note saying it was only test infrastructure.

WHERE THE REMAINING RISK IS

Every defect came from **older hand-built screens** — the booking and job-order flows, the auth funnel, the settings tiles, the shared buttons. The last thirteen screens added to the matrix, all repository-backed and newer, produced **none**. The newer code was written with responsive layout in mind; the older code was not.

Patterns and remedies are recorded in `docs/runbooks/LAYOUT_PATTERNS.md`.

SIX SCREENS REMAIN UNCOVERED, NAMED RATHER THAN OMITTED

`JobOrderScreen` and `ItemOptionsMenuScreen` need a `MerchantServiceModel` with seventeen required fields; `PaymentWebViewScreen` and `CategoryExperienceScreen` need model arguments; `LiveTrackingScreen` needs a `TrackingDataSource` that `main_injector` references but does not register under that name; and `SplashScreen` is blocked on `flutter_secure_storage`'s platform channel, which never completes under the test binding.

Section E — Deploy notes (backend)

Recorded here because the deploy is now manual and undocumented anywhere else. **There is no CI.** Actions are disabled on `Upupapp/servana_api` and on `Upupapp/ServanaClientAPP`.

WHERE PRODUCTION ACTUALLY LIVES

Thing	Path
App directory (PM2 cwd)	<code>/var/www/servana_api</code>
Entry point	<code>/var/www/servana_api/dist/app.js</code> , node, fork mode, no args
PM2 process	<code>servana-prod</code> , port 8000, behind nginx
Environment	<code>/etc/servana_api.env</code> (<code>root:root 600</code>)
Host	192.46.224.126 — 961 MB RAM, 1 CPU
Old location	<code>/home/github-runner/actions-runner/_work/servana_api/servana_api</code> — no longer used , kept as a rollback

THE DEPLOY, IN THREE COMMANDS

```
ssh -n root@192.46.224.126 'cd /var/www/servana_api; npm run build'
ssh -n root@192.46.224.126 'set -a; . /etc/servana_api.env; set +a; pm2 restart servana-prod --update-env'
ssh root@192.46.224.126 'pm2 save'
```

The `set -a` prefix is load-bearing, not decoration. `--update-env` re-reads the environment of the shell PM2 is invoked from — it does not read any `.env` file. And `dotenv` does not rescue it, because the app builds its DB pool before `dotenv.config()` runs (B7). Restart without the prefix and the app comes up with no database password, which is exactly the outage in Section B.

VERIFYING A DEPLOY

```
curl -s https://api.servana.com.ph/readyz      # expect "phase":"ready","ready":true
curl -s https://api.servana.com.ph/api/catalog/summary # expect 3 / 12 / 95
```

The project also ships its own checks: `npm run deploy:verify` and `npm run db:diagnose`. The latter is read-only and is what identified B8.

RULES THAT CAME OUT OF THE 19 AUGUST OUTAGE

- **Never run production from a directory CI can write to.** That single fact turned a failing test into a customer-facing outage.
- **Never run the test typecheck on this box.** It needs more heap than 961 MB allows and aborts at exit 134 (B9).
- **Env files must have no spaces around `=`.** `KEY = value` and `KEY= value` both fail to source, silently, leaving the variable absent. `dotenv` tolerates them, so the file looks fine until something sources it.
- **A running Node process is not evidence that its files exist.** It serves from memory. `/healthz` returning 200 says nothing about whether the next restart will succeed.
- **Check `pm2 env 0`, not the env file,** when a variable seems missing. They disagreed all day.

Section D — How to run this

Nothing is blocked any more. B1 and B2 are closed, the backend is healthy, and TABs 02, 04 and 05 are done. Every remaining TAB can be started today.

1. **TAB 01 first** — fix what “end to end” means before measuring against it. The suite has 2179 tests and none of them walks a customer from install to review.
2. **TAB 03 and TAB 06 next.** Both were blocked by the outage and are the two that touch money and identity: can a new customer actually get in, and can every category actually be booked.
3. **Then TABs 07–11 in order** — they follow the customer's own path, so a break in one is felt by the next.
4. **TAB 12 and 13 last**, because they are about the build rather than the journey.
5. **Re-read Section E before any backend deploy.** The three commands are not interchangeable with a bare `pm2 restart`.

Standing rule for every TAB: finish with evidence — the command run, the status codes seen, the screen reached. A TAB that ends in an assertion has not been done. Where a gate is added, watch it fail once before believing it passes.